

Name: ADARSH REDDY BODHA

Linkedin:www.linkedin.com/in/abodha13reddy | Phno: +1(281) 8101863 | adarshbodha214@gmail.com

SR. BIG DATA ENGINEER

PROFESSIONAL SUMMARY:

- Designed and maintained a **C#/.NET** data integration application utilizing **ADO.NET**, dependency injection, and class inheritance to enable seamless connectivity with **EMR systems**, **Microsoft SQL Server**, and **Azure SQL**, ensuring scalability and healthcare data standards compliance throughout deployment.
- Developed comprehensive **ETL** pipelines within **Azure Databricks** by creating **Databricks notebooks** that extracted clinical data, applied transformations, and loaded processed information into **Microsoft SQL Server** and **Databricks SQL** with validation checkpoints and error handling mechanisms.
- Engineered complex queries using **Microsoft SQL Server** and **Databricks SQL** with stored procedures, views, joins, and indexes to extract and transform large-scale healthcare datasets while maintaining data integrity, security compliance, and optimal performance for clinical applications.
- Collaborated with healthcare providers to validate data accuracy by conducting reconciliation processes, implementing validation rules within **C#/.NET** applications, and documenting data mapping specifications between **EMR systems**, **Microsoft SQL Server**, and **Databricks SQL** databases throughout integration lifecycle.
- Deployed the **C#/.NET** application across healthcare facilities by managing installations, coordinating with IT teams, ensuring **ADO.NET** connectivity to **Microsoft SQL Server**, and implementing security protocols, network configurations, and HIPAA-compliant access controls for clinical data protection.
- Built integration solutions in **Azure** environments by configuring **Databricks** workspaces, managing **Azure SQL** databases, developing **C#/.NET** applications leveraging **Azure services**, and establishing secure connectivity patterns with retry logic for continuous healthcare data processing across distributed systems.
- Developed **CI/CD pipelines** using **Azure DevOps** to automate deployment of **C#/.NET** applications, **Databricks notebooks**, and **Microsoft SQL Server** schema changes across environments, implementing automated testing, code quality checks, and validation procedures for reliable releases with minimal downtime.
- Implemented containerization using **Docker** for **Windows** and **Linux** to package the **C#/.NET** application with dependencies, enabling consistent deployments and orchestration through **Azure Kubernetes Service (AKS)** for scalability, resource optimization, and simplified lifecycle management in cloud environments.
- Managed healthcare **API** integrations by developing **C#/.NET** services consuming RESTful **APIs** from **EMR systems**, implementing authentication, handling rate limiting, configuring **VPN** connections, and troubleshooting connectivity to ensure uninterrupted data flow to **Azure Databricks** and **Microsoft SQL Server**.
- Created **PowerShell** automation scripts for deployment, **Microsoft SQL Server** maintenance, log analysis, and monitoring across healthcare environments, integrating scripts into **CI/CD pipelines** through **Azure DevOps** to reduce manual intervention and ensure consistent **C#/.NET** application configuration.
- Performed comprehensive troubleshooting by analyzing **ETL** workflow failures in **Databricks**, investigating **SQL** query performance in **Microsoft SQL Server**, debugging **C#/.NET** application errors, and resolving **EMR systems** connectivity issues through systematic root cause analysis and preventive measures.
- Designed data validation frameworks within **C#/.NET** applications incorporating **ADO.NET**, dependency injection, and class inheritance to verify clinical data completeness before loading into **Databricks SQL** and **Microsoft SQL Server**, generating comprehensive validation reports for clinical teams review.
- Managed **Azure** infrastructure including **Azure Kubernetes Service (AKS)** cluster administration, **Databricks** workspace configuration, virtual networks, and storage accounts to maintain secure **C#/.NET** application hosting environments with backup strategies, disaster recovery procedures, and monitoring solutions for business continuity.

- Developed data transformation logic in **Databricks notebooks** using **SQL** and **Python** to cleanse and normalize clinical data extracted through **C#/.NET ETL** applications from **EMR systems**, implementing business rules, standardization procedures, and optimization strategies for large patient record volumes.
- Implemented security controls by encrypting data, managing **Microsoft SQL Server** permissions, configuring **Azure** role-based access controls, implementing audit logging in **C#/.NET** applications, and ensuring **ADO.NET** connections utilized secure authentication for HIPAA compliance throughout integration processes.
- Orchestrated **Java**-based microservices within **Azure Kubernetes Service (AKS)** clusters that consumed **APIs** from **EMR systems**, processed clinical data using **Java** frameworks, and published results to **Azure SQL** databases, ensuring seamless interoperability with existing **C#/.NET** integration workflows.
- Deployed **Java** applications using **Docker** containers on **Linux** environments within **Azure**, leveraging **Azure DevOps CI/CD pipelines** to automate builds, implement **ETL** processing logic, and integrate with **Databricks** and **Microsoft SQL Server** for comprehensive healthcare data workflows.
- Developed **Java**-based data processors connecting to **Azure SQL** and **Databricks SQL** platforms using **JDBC**, implementing **API** integrations with **EMR systems**, orchestrating deployments through **Azure Kubernetes Service (AKS)**, and managing configurations via **PowerShell** and **Azure DevOps** automation frameworks.

TECHNICAL SKILLS:

- **ETL:** Informatica Power Center 10.x/9.6, AWS Glue, Talend 5.6, SSIS, Semarchy xDM
- **Databases & Tools:** MS SQL Server 2014/2012, Teradata 15/14, Oracle 11g/10g, Erwin 8/9, MySQL
- **Cloud Environment:** AWS (Snowflake, RDS, Redshift, EC2, EMR, S3, Lambda, Glue, Athena, SQS, VPC, Cloud Watch, Jenkins), Azure (HDInsight, CosmosDB, Fabric, Blob storage, Data lake, Data factory, SQL, Databricks), GCP (DataProc, Bigquery, Cloud storage, GKE, Functions, Spanner)
- **Reporting Tools:** Tableau, PowerBI, EHR, Semarchy
- **Big Data Ecosystem:** HDFS, Hive, MapReduce, Sqoop, Flume, Spark, Spark Streaming, NiFi, Yarn, Zookeeper, Airflow, Kafka
- **Programming languages:** Unix, SQL, Python, T-SQL, Java, Pyspark, Spring boot, Scala
- **Software Life Cycle/ Methodologies:** SDLC, Waterfall, and Agile models
- **Containerization and Orchestration Tools:** Kubernetes, Docker, Jenkins, Docker Hub
- **Python Libraries / Machine Learning Algorithms:** Pandas, NumPy, Matplotlib, Scikit-Learn, Decision Tree, Random Forest, Logistic Regression, Gradient Boosting
- **Data Warehousing & BI:** Star and Snowflake schema, Facts and Dimensions tables, SSIS, Splunk, Cloudera

PROFESSIONAL EXPERIENCE:

Client: Johnson & Johnson

New Brunswick, NJ

Role: SR. BIG DATA ENGINEER

October 2023 – Current

Responsibilities:

- Developed **C#/.NET** applications using **ADO.NET**, **Dependency Injection**, and **Class inheritance** to integrate **EMR** and **Practice Management systems**, ensuring **HIPAA-compliant patient data security** while deploying on **Azure Virtual Machines** and orchestrating **data integration** workflows.
- Designed **ETL pipelines** using **Azure Data Factory**, **Databricks notebooks**, **Apache Spark**, and **Databricks SQL** to extract and transform clinical data from **EMR systems**, applying **HL7 Standards** for interoperability and validating **data quality** before downstream consumption.
- Authored complex **SQL queries** and stored procedures in **Microsoft SQL Server**, **Azure SQL Database**, **Synapse Analytics**, and **Databricks SQL** to perform multi-source aggregation, enabling real-time clinical reporting while maintaining **HITECH** compliance and **patient data security**.

- Configured **Care Path Agent** installations across on-premises and **Azure** environments, establishing **VPN** tunnels, **APIs** for bi-directional exchange with **Practice Management** systems, and troubleshooting **ETL workflows** using **Azure Blob Storage** and infrastructure resources.
- Built **CI/CD pipelines** using **Jenkins**, **Git**, and **PowerShell** to automate deployment of **C#/.NET** applications and **Databricks notebooks**, deploying containerized services on **Azure Kubernetes (AKS)** across **Windows** and **Linux** environments for scalable workflows.
- Collaborated with clinical teams to validate **data quality** of extracted datasets, performing root cause analysis and iterating **ETL** logic in **Databricks SQL** and **Apache Spark** to ensure **HIPAA** compliance and alignment with **HL7 Standards**.
- Developed reusable **C#** libraries using **ADO.NET** and **Dependency Injection** to standardize database access and enable **Class inheritance**-based extensibility, integrating with **Azure Functions** and **APIs** for event-driven **data integration** across **Healthcare applications**.
- Implemented real-time streaming solutions using **Apache Kafka** to ingest clinical events from **EMR systems**, processing through **Apache Spark** on **Databricks** and persisting into **Azure Cosmos DB** and **Snowflake** while ensuring **patient data security**.
- Containerized **C#/.NET** services using **Docker** on **Windows** and **Linux**, deployed to **Azure Kubernetes (AKS)**, and configured monitoring with **Prometheus** and **Grafana** to track performance and ensure uptime of **data integration** services supporting **Clinical applications**.
- Created data models and lineage documentation using **Erwin** and **Collibra** across **Azure SQL Database**, **Synapse Analytics**, and **Snowflake**, enforcing **HIPAA**-compliant data classification and masking policies within **ETL workflows** and **Databricks** environments.
- Executed complex **data transformations** in **Databricks notebooks** using **PySpark** and **Databricks SQL**, applying clinical coding rules and validating output against **HL7 Standards** to ensure accuracy for **Tableau** dashboards and operational reporting systems.
- Managed **Azure** infrastructure including **Virtual Machines**, **Blob Storage**, **Azure SQL Database**, **Azure Functions**, **Cosmos DB**, **Synapse Analytics**, **Data Factory**, and **Azure Monitor**, configuring security groups and **Active Directory** authentication for **HIPAA**-compliant hosting.
- Automated provisioning and deployment using **PowerShell** scripts for **Azure** resources, **Databricks** workspace setup, and **C#/.NET** applications, integrating with **CI/CD pipelines** via **Jenkins** and maintaining version control through **Git** repositories.
- Optimized **SQL** query performance in **Microsoft SQL Server** and **Databricks SQL** through indexing and partitioning, improving **ETL** job efficiency and reducing latency for **APIs** serving clinical data to **Healthcare applications** and **Practice Management systems**.
- Integrated **Azure Monitor** and **Application Insights** with **C#/.NET** applications and **Databricks** workloads to implement logging and alerting, enabling proactive identification of **data integration** failures and tracking incident resolution through **Jira** workflows.
- Designed secure **API** authentication using **Azure Active Directory**, implemented certificate-based **VPN** connections for **EMR** access, and applied encryption across **Azure SQL Database**, **Blob Storage**, and **Cosmos DB** to maintain **HIPAA** compliance and **patient data security**.
- Coordinated user acceptance testing with clinical teams for **ETL** outputs and **Care Path Agent** functionality, resolving data discrepancies and iterating **Databricks notebooks** and **SQL** queries to align with clinical workflows and regulatory requirements.

Environment: Git, Erwin, Tableau, Microsoft SQL, EMR, Windows, Blob Storage, Grafana, C#, Azure Kubernetes, Apache Kafka, Linux, ADO.NET, Jenkins, Data Factory, AKS, Apache Spark, Azure SQL Database, Databricks, Powershell, Azure, Cosmos DB, SQL, Virtual Machines, Databricks Notebooks, .NET, Snowflake, Jira, Synapse Analytics, Prometheus, Collibra, Databricks SQL

Client: eBay

Role: SR DATA ENGINEER

Responsibilities:

San Jose, CA

December 2020 – September 2023

- Designed **C#/.NET** applications using **ADO.NET**, **Dependency Injection**, and **Class Inheritance** to extract transaction data from **Microsoft SQL Server**, transform it through **Databricks**, and load into **Azure SQL Database** while maintaining **PCI-DSS** compliance throughout workflows.

- Developed **SQL** stored procedures in **Microsoft SQL Server** and **Databricks SQL** to integrate order and inventory data, applied **data quality validation** rules, and enabled **Power BI** analytics through **Azure Data Factory** orchestrated pipelines.
- Built **ETL pipelines** using **Azure Data Factory**, **Apache NiFi**, and **Databricks** to ingest data from **Microsoft SQL Server** and **APIs** into **Azure Blob Storage** and **Snowflake** while ensuring **GDPR** and **PCI-DSS** compliance.
- Created reusable **C#** agents with **ADO.NET** and **Class Inheritance**, deployed on **Azure Windows** and **Linux** virtual machines, and monitored through **ELK Stack** for distributed **ETL** execution and operational troubleshooting.
- Authored **Databricks** notebooks in **SQL** and **Python** to execute **data transformations**, perform **data quality validation** including schema drift detection, and integrate results into **Synapse Analytics** for enterprise reporting and segmentation.
- Configured **CI/CD pipelines** using **GitHub Actions** and **Azure DevOps** to automate deployment of **C#/.NET** services and **Databricks** notebooks, managed secrets via **Azure Key Vault**, and maintained **PCI-DSS** audit trails.
- Deployed containerized **C#** microservices using **Docker** on **Azure Kubernetes Service** with **Windows** and **Linux** images, implemented auto-scaling policies, and monitored performance through **Azure Monitor** for distributed **ETL** execution.
- Integrated payment gateway **APIs** using **C#**, **Azure Functions**, and **Java** microservices with **VPN** tunneling and **OAuth** authentication, enforcing **PCI-DSS** tokenization standards and role-based access control for secure **data integration**.
- Developed **PowerShell** scripts to provision **Azure** resources, schedule **SQL** jobs on **Microsoft SQL Server**, and orchestrate **Databricks**, **Azure Data Factory**, and **Snowflake** dependencies for automated nightly **ETL** batch processing.
- Performed **data quality validation** by defining rules with business teams, auditing discrepancies using **SQL** and **Databricks** aggregations, and documenting workflows in **Confluence** to maintain **GDPR** and **PCI-DSS** compliance standards.
- Administered **Azure SQL Database**, **Cosmos DB**, and **Synapse Analytics** by implementing backup strategies, indexing optimizations, and query tuning using **Microsoft SQL** tools while ensuring **data transformations** met regulatory security requirements.
- Utilized **dbt** to define version-controlled **SQL** transformation models in **Snowflake** and **Databricks**, enabled dimensional modeling for analytics, and tracked lineage in **Alation** for auditable **data quality validation** workflows.
- Configured **Azure Blob Storage** with lifecycle policies, encryption, and **GDPR/PCI-DSS** compliant access controls, orchestrated downstream **ETL** through **Azure Data Factory** triggers and **Databricks** job clusters for optimized performance.
- Designed **Class Inheritance** hierarchies in **C#** for extensible connectors supporting **Microsoft SQL Server**, **APIs**, and message queues, applied **Dependency Injection** for configuration, and enabled rapid platform onboarding without duplication.
- Integrated **ER/Studio** for data modeling and metadata documentation, maintained entity-relationship diagrams and lineage mappings for **SQL** databases and **Databricks** tables, and facilitated impact analysis during **ETL** enhancements.
- Monitored **data integration** workflows using **Azure Monitor**, **ELK Stack**, and **Databricks** logs to resolve **VPN**, **API**, and database issues, implemented retry logic and circuit breakers in **C#** and **Java** applications.

Environments: ETL, AKS, er/studio, Containerization, APIs, github actions, Azure SQL Database, elk stack, confluence, snowflake, apache nifi, power bi, Databricks, Functions, CI/CD pipelines, Cosmos DB, VPN, Databricks notebooks, Microsoft SQL, ADO.NET, SQL, Synapse Analytics, .NET, Blob Storage, Azure Kubernetes, Azure, Virtual Machines, GitHub, Monitor, PowerShell, Data Factory, alation, dbt, Windows, Linux, C#

Client: AT & T

Role: DATA ENGINEER

Responsibilities:

Dallas, TX

February 2018 – November 2020

- Developed **C#/.NET** data integration applications using **ADO.NET**, **Dependency Injection**, and **Class Inheritance** to extract telecom subscriber and billing data from legacy OSS/BSS systems, transforming and loading into **Microsoft SQL Server** and **Azure SQL Database** with **Great Expectations** validation frameworks for data quality assurance.
- Built **ETL pipelines** using **Azure Data Factory**, **Databricks**, and **Apache Beam** to orchestrate transformations across CDR logs and network traffic data, writing **SQL** queries and **Databricks SQL** notebooks to cleanse and partition datasets for regulatory reporting and downstream analytics consumption.

- Designed **data integration** workflows in **C#** leveraging **ADO.NET** connection pooling and **Class Inheritance** for telecom source connectors, ensuring compatibility with **Azure Blob Storage**, **Azure SQL Database**, **Cosmos DB**, and **Snowflake** while maintaining schema versioning and backward compatibility across consumers.
- Created reusable **Databricks** notebooks in **Python** and **SQL** for incremental transformations applying telecom KPI business logic, integrating with **Azure Synapse Analytics** for star schema modeling and publishing curated datasets to **Looker** dashboards for executive teams and operational reporting.
- Implemented **CI/CD pipelines** using **GitLab CI** to automate build, test, and deployment of **C#/.NET** agents and **Databricks** notebooks across environments on **Azure Kubernetes Service**, incorporating **containerization** with **Docker** images for consistent runtime environments and dependency isolation.
- Configured **Azure Kubernetes Service** clusters to deploy **containerized C#** extraction agents pulling network event data via secure **APIs** and **VPN** tunnels, managing pod scaling and persistent volume claims linked to **Azure Blob Storage** for staging raw data files on scheduled intervals.
- Wrote optimized **SQL** queries and stored procedures in **Microsoft SQL Server** using joins and window functions for telecom billing reconciliation, implementing indexing strategies and integrating **data quality validation** checks with **Splunk** logging for real-time anomaly detection and alerting.
- Developed **PowerShell** scripts to automate infrastructure provisioning and deployment of **C#** integration agents on **Windows** VMs and **AKS** nodes, including secret rotation and **Azure Monitor** integration for health checks, reducing deployment time and manual intervention significantly.
- Designed **Azure Functions** in **C#** using **Dependency Injection** and **ADO.NET** to trigger event-driven **ETL** workflows upon new file arrivals in **Azure Blob Storage**, orchestrating downstream processing via **Azure Data Factory** and logging execution metadata to **Azure SQL Database**.
- Integrated **Flume** agents and **Apache Beam** streaming pipelines to capture real-time telecom CDR streams, transforming JSON and XML payloads into normalized schemas stored in **Microsoft SQL Server** and **Databricks** Delta Lake with custom **C#** UDFs for telecom-specific parsing logic.
- Collaborated with operations and finance teams to validate transformed datasets using **Great Expectations** profiling suites, defining custom expectation rules for subscriber demographics and billing accuracy, publishing validation reports to **Slack** channels and **Looker** for stakeholder review and approval.
- Deployed **Databricks** workspaces on **Azure** with cluster policies and RBAC controls, integrating **Azure Active Directory** for secure access to **Databricks** notebooks and **Databricks SQL** warehouses, enabling self-service analytics while maintaining centralized governance and comprehensive lineage tracking.
- Built reusable **C#** libraries implementing **Class Inheritance** and interface patterns to standardize extraction from Oracle databases, flat files, REST **APIs**, and SFTP endpoints, encapsulating connection logic and retry mechanisms within base classes extended per source-specific requirements.
- Established monitoring infrastructure using **Splunk**, **Azure Monitor**, and **Azure Application Insights** to capture runtime metrics and exceptions from **C#/.NET** agents and **Databricks** ETL jobs, configuring alerts for SLA breaches and **data quality** threshold violations for proactive incident response.
- Managed **Azure SQL Database** elastic pools applying query tuning, index optimization, and partitioning strategies for large telecom transaction tables, ensuring sub-second query response times for concurrent read/write operations from **ETL** workflows and analytical reporting workloads supporting business operations.
- Automated deployment of **containerized C#** agents using **GitLab CI** and **PowerShell** scripts targeting **AKS** clusters with multi-stage builds and vulnerability scanning, ensuring zero-downtime deployments and rapid rollback capabilities while maintaining secure **VPN** connectivity for sensitive telecom data integration.

Environments: Azure Kubernetes, .NET, flume, looker, APIs, ADO.NET, C#, Cosmos DB, GitLab, Blob Storage, Azure, great expectations, Windows, apache beam, Linux, CI/CD pipelines, Databricks notebooks, Microsoft SQL, Functions, Synapse Analytics, splunk, slack, SQL, power designer, Virtual Machines, PowerShell, snowflake, gitlab ci, Databricks, Data Factory, Azure SQL Database, AKS

Client: Honeywell

Role: DATA ENGINEER

Responsibilities:

Charlotte, NC

December 2015 – January 2018

- Developed comprehensive Manufacturing Data Integration platform using **Azure Cloud**, standardizing data flows between production systems and business intelligence tools. Improved data integration efficiency by 65% while reducing manual reconciliation efforts across global manufacturing facilities.
- Created data ingestion pipelines through **Talend ETL** and **Azure Data Factory (ADF)** for processing industrial sensor data from production lines and quality control systems. Implemented automated validation checks ensuring data accuracy from over 200 manufacturing data sources worldwide.
- Executed data migration projects transferring production records and equipment specifications from legacy mainframes to modern manufacturing execution systems. Successfully transferred historical manufacturing data with 100% integrity validation and zero production downtime.
- Built complex workflows with **Apache Airflow**, coordinating ETL processes between shop floor control systems and enterprise resource planning platforms. Implemented intelligent scheduling algorithms reducing processing bottlenecks during shift changes and production peaks.
- Set up **Azure SQL Database** for structured manufacturing data storage, implementing advanced indexing strategies for production metrics and quality control records. Optimized query performance reducing reporting response times by 78% for critical production monitoring dashboards.
- Implemented **Azure Data Lake Storage (ADLS)** for scalable storage of semi-structured equipment telemetry data and unstructured maintenance documentation. Created hierarchical namespace organization enabling efficient access to historical production data while maintaining regulatory compliance.
- Built data security protocols meeting **ISO 27001** and industrial regulatory requirements while working with proprietary manufacturing data models and industry standards. Implemented comprehensive encryption and access controls protecting intellectual property across global manufacturing operations.
- Executed migration of legacy manufacturing control systems to **SAP**, ensuring seamless integration with production equipment and compliance with industrial protocols. Developed custom connectors for specialized manufacturing equipment ensuring continuous data collection during migration phases.
- Created predictive models using **Azure Databricks**, analyzing equipment performance patterns to forecast maintenance needs before failures occurred. Improved production line uptime by 23% through early identification of potential equipment issues using machine learning techniques.
- Built interactive **Tableau** dashboards visualizing production metrics and quality indicators across global manufacturing facilities. Enabled real-time decision-making through customized executive views highlighting production efficiency and compliance metrics.
- Designed real-time monitoring system for industrial equipment using **IoT Hub**, processing thousands of sensor readings per second from production lines. Implemented anomaly detection algorithms triggering immediate alerts when operating conditions deviated from optimal parameters.
- Developed custom data transformation frameworks in **Python** for normalizing inconsistent manufacturing specifications from acquired companies. Created standardized reference data models ensuring consistent quality reporting across diverse product lines and manufacturing facilities.
- Implemented **Azure Synapse Analytics** for enterprise-wide manufacturing intelligence, consolidating production metrics from multiple global facilities. Created dimensional data models optimized for complex analytical queries across product lines, regions, and time periods.

Environments: Azure Cloud, Talend ETL, Azure Data Factory (ADF), Azure Databricks, Azure Data Lake Storage, Tableau, Epic, FHIR.

Client: Ramco Systems

Role: ETL TESTER

Responsibilities:

Chennai, India

May 2014 – September 2015

- Developed scalable ETL pipeline using **AWS**, **Informatica**, and **Talend** integrating data from aviation maintenance systems and enterprise resource planning platforms. Successfully consolidated information from multiple aircraft maintenance facilities, improving data consistency and reporting accuracy for global airline clients.
- Set up **AWS S3** storage for secure management of raw, staging, and transformed aircraft maintenance records and parts inventory data. Created efficient data ingestion workflows reducing processing time for critical maintenance scheduling information by 40% while maintaining regulatory compliance.
- Used **Informatica PowerCenter** extracting data from **MySQL**, **PostgreSQL**, and **MongoDB** databases containing aircraft service records and component lifecycle information. Ensured precise data extraction with custom transformations handling complex relationships between maintenance events and aircraft configurations.
- Created **Talend** jobs processing CSV and JSON file formats containing flight operations data and maintenance logs from diverse airline systems. Successfully reduced data processing time for regulatory reporting by implementing parallel processing strategies and optimized transformation logic.
- Built detailed data transformations with **Informatica**, focusing on data cleansing and quality verification of aircraft maintenance records and regulatory compliance documentation. Increased accuracy of maintenance forecasting by implementing sophisticated validation rules detecting anomalies in component service history.
- Set up **Talend Open Studio** for data transformations and integrations between legacy aviation systems and modern cloud platforms. Made ETL workflows efficient through reusable components and standardized transformation patterns, reducing development time by 20% for new airline client onboarding.
- Developed **AWS Glue** jobs to catalog and move transformed aircraft maintenance and inventory data into **Redshift** for analytics purposes. Improved data storage and retrieval performance enabling maintenance planners to quickly access historical service records across entire aircraft fleets.
- Built **Tableau** connection to **AWS Redshift**, creating interactive dashboards visualizing aircraft maintenance schedules and component reliability metrics. Enabled real-time data analysis and business insights helping maintenance operations teams optimize resource allocation and improve aircraft availability.

Environments: AWS, Informatica, Talend, AWS S3, AWS Glue, AWS Redshift, Tableau, AWS Lambda.

EDUCATION:

Degree: Bachelor of Technology in Computer Science and Engineering (May 2010 – March 2014)

College: MRCET

GPA: 7.81