

Rafael A. Mendes

VP DATA ENGINEERING / CHIEF DATA OFFICER

Email: rafael.mendes@data-architect.pro

Phone: +1 (512) 904-7732

Location: Austin, TX 78701

LinkedIn: linkedin.com/in/rafael-a-mendes

15+
Yrs Data Eng

500TB+
Data Managed

\$220M
Cost Savings

8
Data Platforms Built

40+
Data Engineers Led

PROFESSIONAL SUMMARY

Visionary data engineering leader with 15+ years architecting and scaling enterprise data platforms, data lakes, and real-time analytics systems. Track record building 8 production data platforms handling 500TB+ daily data at Dell Technologies, Netflix, Uber, and Cloudera. Expert in lakehouse architecture, streaming data pipelines, data governance, and modern data stack. Deep hands-on expertise in Apache Spark, Kafka, Flink, dbt, and Snowflake. Passionate about data democratization, self-serve analytics, and building data engineering cultures.

SKILLS

DATA ARCHITECTURE

Lakehouse Data Mesh
Lambda/Kappa
Data Warehouse
Data Lake Data Vault

PROCESSING ENGINES

Apache Spark
Apache Flink
Kafka Streams
Apache Beam Dask Ray

STREAMING

Apache Kafka
Amazon Kinesis
Google Pub/Sub
Apache Pulsar
Confluent Platform

STORAGE & QUERY

Snowflake Databricks
BigQuery Redshift
Delta Lake
Apache Iceberg dbt

ORCHESTRATION

Apache Airflow Prefect
Dagster Luigi Temporal

GOVERNANCE & QUALITY

Apache Atlas Alation
Great Expectations
OpenMetadata
Monte Carlo

LANGUAGES

Python (Expert)
Scala (Advanced)
Java (Proficient)
SQL (Expert) Bash

CERTIFICATIONS

Databricks Certified Data Engineer Professional (2022)

AWS Data Analytics Specialty (2022)

Google Professional Data Engineer (2021)

Snowflake Advanced Architect

WORK EXPERIENCE

VP Data Engineering Dell Technologies · Austin, TX

Feb 2021 - Present

- Lead 45-person data engineering organization supporting Dell's \$92B global revenue operations; oversee 8 mission-critical data platforms handling 580TB daily
- Architected Dell's enterprise data lakehouse on Databricks + Delta Lake + AWS S3; migrated from 3 legacy Hadoop clusters, achieving 68% query performance improvement and \$18M annual cost reduction
- Built real-time supply chain analytics platform processing 420M events/day from 12,000+ supplier systems; reduced demand forecasting error from 18% to 4.2%, avoiding \$140M in excess inventory
- Designed data mesh architecture with 14 domain teams; increased data product availability from 72% to 99.2% through federated ownership model
- Led data governance program: implemented data catalog (Alation), lineage (OpenLineage), quality monitoring (Great Expectations) for 50,000+ datasets
- Established ML Feature Store enabling 120 data scientists to share and reuse 3,400+ features; reduced model development cycle from 8 weeks to 10 days

Databricks Delta Lake Spark Kafka dbt Airflow Snowflake Python Scala
AWS

Senior Data Engineering Manager Netflix · Los Gatos, CA

Mar 2017 - Jan 2021

- Led 22-person Data Platform team building streaming analytics infrastructure for Netflix's 240M+ subscribers across 190 countries
- Co-designed Netflix's next-gen data lake on Apache Iceberg (project co-founder); Iceberg has since become Apache Top-Level Project with 800+ contributors
- Built real-time content popularity tracking system processing 150B events/day; powered A/B testing, recommendation ranking, and thumbnail optimization experiments
- Launched self-serve data platform enabling 4,000+ Netflix employees to run ad-hoc queries without engineering dependency; reduced data request SLA from 2 weeks to 2 hours

Apache Iceberg Spark Kafka Flink Presto Hive Python Scala AWS

Staff Data Engineer Uber Technologies · San Francisco, CA

Jan 2013 - Feb 2017

- Early data engineer at Uber (#12 in data org); built analytics infrastructure from 0 as company scaled from 5 to 75 cities and \$1B to \$5B GMV
- Designed Uber's first real-time surge pricing data pipeline; processed 40M GPS pings/minute with <200ms latency using Kafka + Storm
- Built Uber's internal data warehouse from scratch on Vertica, later migrated to Hive on S3; supported 200+ analysts and data scientists

Kafka Storm Hive Vertica Python Java MySQL

Data Engineer / ETL Developer Cloudera, Inc. · San Jose, CA

Jun 2009 - Dec 2012

- Contributed to Cloudera's early Hadoop ecosystem tooling; co-developed initial version of Sqoop (Hadoop-RDBMS connector), now an Apache top-level project with 500+ contributors
- Consulted with 25+ enterprise customers on Hadoop adoption strategy; implemented data lake pilots at 8 Fortune 500 companies

Hadoop Hive Pig Sqoop Java Python

(2021)

Confluent Certified Kafka Developer (2020)

dbt Analytics Engineering Cert (2022)

EDUCATION

M.S., Computer Science
University of Texas at Austin
2009
Specialization: Data Systems · GPA 3.95

B.S., Information Systems
Federal University of São Paulo (UNIFESP)
2007
Valedictorian · Full Merit Scholarship

LANGUAGES

English	Fluent (C2)
Portuguese	Native
Spanish	Fluent (C1)
French	Basic (A2)

AWARDS

- 🏆 Databricks Community Champion (2022)
- 🏆 Apache Software Foundation Committer – Iceberg, Sqoop
- 🏆 Austin Data Council Speaker of the Year (2021)
- 🏆 Dell Technologies Data Excellence Award (2023)

VOLUNTEERING

- Apache Iceberg PMC Member (2020–Present)
- Strata Data Conference Program Committee (2019–Present)
- Data Engineering Podcast – Monthly Guest Expert
- Mentor: LatinX Data Scientists Network

CONTACT

✉ rafael.mendes@data-architect.pro

☎ +1 (512) 904-7732

△ 300 W 6th Street, Austin, TX 78701

KEY PROJECTS

Dell Enterprise Lakehouse
2021–2022

Ground-up enterprise data lakehouse on Databricks+Delta Lake. 580TB daily data, 68% faster queries, \$18M annual savings. Replaced 3 Hadoop clusters and 2 legacy EDWs.

Apache Iceberg (Co-founder)
2017–2019

Co-designed Apache Iceberg table format at Netflix. Now Apache Top-Level Project, used by Airbnb, Lyft, LinkedIn, Apple, and 500+ orgs. 6,000+ GitHub stars.

PUBLICATIONS & TALKS

- ▶ Apache Iceberg: The Definitive Guide – O'Reilly, 2023 (Co-author)
- ▶ Strata Data Conference Keynote – Modern Data Architecture (2022)
- ▶ Data Engineering Podcast Episodes 180, 230 (2021, 2022)
- ▶ Apache Iceberg Technical Blog – 8 posts, 200K+ reads
- ▶ Netflix Tech Blog: 'Iceberg: a fast table format for S3' (2018, 120K reads)

PARSER VERIFICATION TABLE

Field	Value	Status
name	Rafael A. Mendes	✓ HIGH
email	rafael.mendes@data-architect.pro	✓ HIGH
phone	+1 (512) 904-7732	✓ HIGH
location/address	300 W 6th Street, Austin, TX 78701	✓ HIGH
job_title	VP Data Engineering / Chief Data Officer	✓ HIGH
current_employer	Dell Technologies	✓ HIGH
years_experience	15+	✓ HIGH
work_history_entries	4	✓ HIGH
education_entries	2	✓ HIGH
certifications_count	6	✓ HIGH
skills_groups_count	7	✓ HIGH
languages	English:Fluent (C2), Portuguese:Native, Spanish:Fluent (C1), French:Basic (A2)	✓ HIGH
awards_count	4	✓ HIGH
projects_count	2	✓ HIGH